



Introduction to Statistical Machine Learning

Christfried Webers

Statistical Machine Learning Group
NICTA
and
College of Engineering and Computer Science
The Australian National University

Canberra
February – June 2013

(Many figures from C. M. Bishop, "Pattern Recognition and Machine Learning")

Outlines

Overview
Introduction
Linear Algebra
Probability
Linear Regression 1
Linear Regression 2
Linear Classification 1
Linear Classification 2
Neural Networks 1
Neural Networks 2
Kernel Methods
Sparse Kernel Methods
Graphical Models 1
Graphical Models 2
Graphical Models 3
Mixture Models and EM 1
Mixture Models and EM 2
Approximate Inference
Sampling
Principal Component Analysis
Sequential Data 1
Sequential Data 2
Combining Models
Selected Topics
Discussion and Summary



Part VII

Linear Classification 1

Classification

*Generalised Linear
Model*

Inference and Decision

Discriminant Functions

*Fisher's Linear
Discriminant*

*The Perceptron
Algorithm*



- Goal : Given input data \mathbf{x} , assign it to one of K discrete classes \mathcal{C}_k where $k = 1, \dots, K$.
- Divide the input space into different regions.




Figure : Length of the petal [in cm] for a given sepal [cm] for iris flowers (Iris Setosa, Iris Versicolor, Iris Virginica).

How to represent binary class labels?



- Class labels are no longer real values as in regression, but a discrete set.
- Two classes : $t \in \{0, 1\}$
($t = 1$ represents class \mathcal{C}_1 and $t = 0$ represents class \mathcal{C}_2)
- Can interpret the value of t as the probability of class \mathcal{C}_1 , with only two values possible for the probability, 0 or 1.
- Note: Other conventions to map classes into integers possible, check the setup.

Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

How to represent multi-class labels?



Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- If there are more than two classes ($K > 2$), we call it a multi-class setup.
- Often used: 1-of- K coding scheme in which \mathbf{t} is a vector of length K which has all values 0 except for $t_j = 1$, where j comes from the membership in class C_j to encode.
- Example: Given 5 classes, $\{C_1, \dots, C_5\}$. Membership in class C_2 will be encoded as the target vector

$$\mathbf{t} = (0, 1, 0, 0, 0)^T$$


- Note: Other conventions to map multi-classes into integers possible, check the setup.



- Idea: Use again a **Linear Model** as in regression: $y(\mathbf{x}, \mathbf{w})$ is a linear function of the parameters \mathbf{w}

$$y(\mathbf{x}_n, \mathbf{w}) = \mathbf{w}^T \phi(\mathbf{x}_n)$$

- But generally $y(\mathbf{x}_n, \mathbf{w}) \in \mathbb{R}$.
Example: Which class is $y(\mathbf{x}, \mathbf{w}) = 0.71623$?



Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm



Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Apply a mapping $f : \mathbb{R} \rightarrow \mathbb{Z}$ to the linear model to get the discrete class labels.
- Generalised Linear Model

$$y(\mathbf{x}_n, \mathbf{w}) = f(\mathbf{w}^T \phi(\mathbf{x}_n))$$

- Activation function: $f(\cdot)$
- Link function : $f^{-1}(\cdot)$




Figure : Example of an activation function $f(z) = \text{sign}(z)$.

Three Models for Decision Problems



In increasing order of complexity

- Find a **discriminant function** $f(\mathbf{x})$ which maps each input directly onto a class label.
- Discriminative Models
 - 1 Solve the inference problem of determining the posterior class probabilities $p(\mathcal{C}_k | \mathbf{x})$.
 - 2 Use decision theory to assign each new \mathbf{x} to one of the classes.
- Generative Models
 - 1 Solve the inference problem of determining the class-conditional probabilities $p(\mathbf{x} | \mathcal{C}_k)$.
 - 2 Also, infer the prior class probabilities $p(\mathcal{C}_k)$.
 - 3 Use Bayes' theorem to find the posterior $p(\mathcal{C}_k | \mathbf{x})$.
 - 4 Alternatively, model the joint distribution $p(\mathbf{x}, \mathcal{C}_k)$ directly.
 - 5 Use decision theory to assign each new \mathbf{x} to one of the classes.

Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm



Definition

A **discriminant** is a function that maps from an input vector \mathbf{x} to one of K classes, denoted by \mathcal{C}_k .

- Consider first two classes ($K = 2$).
- Construct a linear function of the inputs \mathbf{x}

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

such that \mathbf{x} being assigned to class \mathcal{C}_1 if $y(\mathbf{x}) \geq 0$, and to class \mathcal{C}_2 otherwise.

- **weight vector** \mathbf{w}
- **bias** w_0 (sometimes $-w_0$ called **threshold**)




- Decision boundary $y(\mathbf{x}) = 0$ is a $(D - 1)$ -dimensional hyperplane in a D -dimensional input space (**decision surface**).
- \mathbf{w} is orthogonal to any vector lying in the decision surface.
- Proof: Assume \mathbf{x}_A and \mathbf{x}_B are two points lying in the decision surface. Then,

$$0 = y(\mathbf{x}_A) - y(\mathbf{x}_B) = \mathbf{w}^T(\mathbf{x}_A - \mathbf{x}_B)$$



- The normal distance from the origin to the decision surface is

$$\frac{\mathbf{w}^T \mathbf{x}}{\|\mathbf{w}\|} = -\frac{w_0}{\|\mathbf{w}\|}$$





Classification

Generalised Linear
Model

Inference and Decision


Discriminant Functions


Fisher's Linear
Discriminant

The Perceptron
Algorithm

- The value of $y(\mathbf{x})$ gives a signed measure of the perpendicular distance r of the point \mathbf{x} from the decision surface, $r = y(\mathbf{x})/\|\mathbf{w}\|$.

$$y(\mathbf{x}) = \mathbf{w}^T \left(\overbrace{\mathbf{x}_\perp + r \frac{\mathbf{w}}{\|\mathbf{w}\|}}^{\mathbf{x}} \right) + w_0 = r \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} + \overbrace{\mathbf{w}^T \mathbf{x}_\perp + w_0}^0 = r \|\mathbf{w}\|$$





Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm


- More compact notation : Add an extra dimension to the input space and set the value to $x_0 = 1$.
- Also define $\tilde{\mathbf{w}} = (w_0, \mathbf{w})$ and $\tilde{\mathbf{x}} = (1, \mathbf{x})$

$$y(\mathbf{x}) = \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}$$

- Decision surface is now a D -dimensional hyperplane in a $D + 1$ -dimensional expanded input space.



- Number of classes $K > 2$
- Can we combine a number of two-class discriminant functions using $K - 1$ **one-versus-the-rest** classifiers?



Classification

Generalised Linear
Model

Inference and Decision


Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm



- Number of classes $K > 2$
- Can we combine a number of two-class discriminant functions using $K(K-1)/2$ **one-versus-one** classifiers?



Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant


The Perceptron
Algorithm



- Number of classes $K > 2$
- Solution: Use K linear functions

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0}$$

- Assign input \mathbf{x} to class \mathcal{C}_k if $y_k(\mathbf{x}) > y_j(\mathbf{x})$ for all $j \neq k$.
- Decision boundary between class \mathcal{C}_k and \mathcal{C}_j given by $y_k(\mathbf{x}) = y_j(\mathbf{x})$



Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Least Squares for Classification



- Regression with a linear function of the model parameters and minimisation of sum-of-squares error function resulted in a closed-form solution for the parameter values.
- Is this also possible for classification?
- Given input data \mathbf{x} belonging to one of K classes \mathcal{C}_k .
- Use 1-of- K binary coding scheme.
- Each class is described by its own linear model

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0} \quad k = 1, \dots, K$$

Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm



- With the conventions

$$\tilde{\mathbf{w}}_k = \begin{bmatrix} w_{k0} \\ \mathbf{w}_k \end{bmatrix} \in \mathbb{R}^{D+1}$$

$$\tilde{\mathbf{x}} = \begin{bmatrix} 1 \\ \mathbf{x} \end{bmatrix} \in \mathbb{R}^{D+1}$$

$$\tilde{\mathbf{W}} = [\tilde{\mathbf{w}}_1 \quad \dots \quad \tilde{\mathbf{w}}_K] \in \mathbb{R}^{(D+1) \times K}$$

- we get for the discriminant function (vector valued)

$$\mathbf{y}(\mathbf{x}) = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}} \in \mathbb{R}^K.$$

- For a new input \mathbf{x} , the class is then defined by the index of the largest value in the row vector $\mathbf{y}(\mathbf{x})$



- Given a training set $\{\mathbf{x}_n, \mathbf{t}\}$ where $n = 1, \dots, N$, and \mathbf{t} is the class in the 1-of- K coding scheme.
- Define a matrix \mathbf{T} where row n corresponds to \mathbf{t}_n^T .
- The sum-of-squares error can now be written as

$$E_D(\tilde{\mathbf{W}}) = \frac{1}{2} \text{tr} \left\{ (\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T})^T (\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T}) \right\}$$

- The minimum of $E_D(\tilde{\mathbf{W}})$ will be reached for

$$\tilde{\mathbf{W}} = (\tilde{\mathbf{X}}^T \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^T \mathbf{T} = \tilde{\mathbf{X}}^\dagger \mathbf{T}$$

where $\tilde{\mathbf{X}}^\dagger$ is the pseudo-inverse of $\tilde{\mathbf{X}}$.



- The discriminant function $\mathbf{y}(\mathbf{x})$ is therefore

$$\mathbf{y}(\mathbf{x}) = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}} = \mathbf{T}^T (\tilde{\mathbf{X}}^\dagger)^T \tilde{\mathbf{x}},$$

where $\tilde{\mathbf{X}}$ is given by the training data, and $\tilde{\mathbf{x}}$ is the new input.

- Interesting property: If for every \mathbf{t}_n the same linear constraint $\mathbf{a}^T \mathbf{t}_n + b = 0$ holds, then the prediction $\mathbf{y}(\mathbf{x})$ will also obey the same constraint



$$\mathbf{a}^T \mathbf{y}(\mathbf{x}) + b = 0.$$

- For the 1-of- K coding scheme, the sum of all components in \mathbf{t}_n is one, and therefore all components of $\mathbf{y}(\mathbf{x})$ will sum to one. BUT: the components are not probabilities, as they are not constraint to the interval $(0, 1)$.

Deficiencies of the Least Squares Approach



Magenta curve : Decision Boundary for the least squares approach (Green curve : Decision boundary for the logistic regression model described later)



Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions



Fisher's Linear
Discriminant

The Perceptron
Algorithm

Deficiencies of the Least Squares Approach



Magenta curve : Decision Boundary for the least squares approach (Green curve : Decision boundary for the logistic regression model described later)



Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm



- View linear classification as dimensionality reduction.

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$$

If $y \geq -w_0$ then class \mathcal{C}_1 , otherwise \mathcal{C}_2 .

- But there are many projections from a D -dimensional input space onto one dimension.
- Projection always means loss of information.
- For classification we want to preserve the class separation in one dimension.
- Can we find a projection which maximally preserves the class separation ?

Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions



Fisher's Linear
Discriminant

The Perceptron
Algorithm

Fisher's Linear Discriminant



Samples from two classes in a two-dimensional input space and their histogram when projected to two different one-dimensional spaces.



Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Fisher's Linear Discriminant - First Try




- Given N_1 input data of class \mathcal{C}_1 , and N_2 input data of class \mathcal{C}_2 , calculate the centres of the two classes

$$\mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n, \quad \mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n$$

- Choose \mathbf{w} so as to maximise the projection of the class means onto \mathbf{w}

$$m_1 - m_2 = \mathbf{w}^T (\mathbf{m}_1 - \mathbf{m}_2)$$

- Problem with non-uniform covariance



Fisher's Linear Discriminant




- Measure also the within-class variance for each class

$$s_k^2 = \sum_{n \in \mathcal{C}_k} (y_n - m_k)^2$$

where $y_n = \mathbf{w}^T \mathbf{x}_n$.

- Maximise the **Fisher criterion**

$$J(\mathbf{w}) = \frac{(m_2 - m_1)^2}{s_1^2 + s_2^2}$$





- The Fisher criterion can be rewritten as

$$J(\mathbf{w}) = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

- \mathbf{S}_B is the **between-class** covariance

$$\mathbf{S}_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T$$

- \mathbf{S}_W is the **within-class** covariance

$$\mathbf{S}_W = \sum_{n \in \mathcal{C}_1} (\mathbf{x}_n - \mathbf{m}_1)(\mathbf{x}_n - \mathbf{m}_1)^T + \sum_{n \in \mathcal{C}_2} (\mathbf{x}_n - \mathbf{m}_2)(\mathbf{x}_n - \mathbf{m}_2)^T$$



- The Fisher criterion

$$J(\mathbf{w}) = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

has a maximum for **Fisher's linear discriminant**

$$\mathbf{w} \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$

- Fisher's linear discriminant is NOT a discriminant, but can be used to construct one by choosing a threshold y_0 in the projection space.

Fisher's Discriminant For Multi-Class



- Assume that the dimensionality of the input space D is greater than the number of classes K .
- Use $D' > 1$ linear 'features' $y_k = \mathbf{w}^T \mathbf{x}$ and write everything in vector form (no bias involved!)

$$\mathbf{y} = \mathbf{W}^T \mathbf{x}.$$

- The within-class covariance is then the sum of the covariances for all K classes

$$\mathbf{S}_W = \sum_{k=1}^K \mathbf{S}_k$$

where

$$\mathbf{S}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T$$
$$\mathbf{m}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{x}_n$$



- Between-class covariance

$$\mathbf{S}_B = \sum_{k=1}^K N_k (\mathbf{m}_k - \mathbf{m})(\mathbf{m}_k - \mathbf{m})^T.$$

where \mathbf{m} is the total mean of the input data

$$\mathbf{m} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n.$$

- One possible way to define a function of \mathbf{W} which is large when the between-class covariance is large and the within-class covariance is small is given by

$$J(\mathbf{W}) = \text{tr} \{ (\mathbf{W}^T \mathbf{S}_W \mathbf{W})^{-1} (\mathbf{W}^T \mathbf{S}_B \mathbf{W}) \}$$

- The maximum of $J(\mathbf{W})$ is determined by the D' eigenvectors of $\mathbf{S}_W^{-1} \mathbf{S}_B$ with the largest eigenvalues.

Fisher's Discriminant For Multi-Class



- How many linear 'features' can one find with this method?
- S_B is of rank at most $K - 1$ because of the sum of K rank one matrices and the global constraint via \mathbf{m} .
- Projection onto the subspace spanned by S_B can not have more than $K - 1$ linear features.

Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

The Perceptron Algorithm



- Frank Rosenblatt (1928 - 1969)
- "Principles of neurodynamics: Perceptrons and the theory of brain mechanisms" (Spartan Books, 1962)



Classification

*Generalised Linear
Model*

Inference and Decision

Discriminant Functions

*Fisher's Linear
Discriminant*

*The Perceptron
Algorithm*

The Perceptron Algorithm



- Perceptron ("MARK 1") was the first computer which could learn new skills by trial and error



Classification

*Generalised Linear
Model*

Inference and Decision

Discriminant Functions

*Fisher's Linear
Discriminant*

*The Perceptron
Algorithm*

The Perceptron Algorithm



- Two class model
- Create feature vector $\phi(\mathbf{x})$ by a fixed nonlinear transformation of the input \mathbf{x} .
- Generalised linear model

$$y(\mathbf{x}) = f(\mathbf{w}^T \phi(\mathbf{x}))$$

with $\phi(\mathbf{x})$ containing some bias element $\phi_0(\mathbf{x}) = 1$.

- nonlinear **activation** function

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

- Target coding for perceptron

$$t = \begin{cases} +1, & \text{if } C_1 \\ -1, & \text{if } C_2 \end{cases}$$

The Perceptron Algorithm - Error Function



Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Idea : Minimise total number of misclassified patterns.
- Problem : As a function of \mathbf{w} , this is piecewise constant and therefore the gradient is zero almost everywhere.
- Better idea: Using the $(-1, +1)$ target coding scheme, we want all patterns to satisfy $\mathbf{w}^T \phi(\mathbf{x}_n) t_n > 0$.
- **Perceptron Criterion** : Add the errors for all patterns belonging to the set of misclassified patterns \mathcal{M}

$$E_P(\mathbf{w}) = - \sum_{n \in \mathcal{M}} \mathbf{w}^T \phi(\mathbf{x}_n) t_n$$



Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Perceptron Criterion (with notation $\phi_n = \phi(\mathbf{x}_n)$)

$$E_P(\mathbf{w}) = - \sum_{n \in \mathcal{M}} \mathbf{w}^T \phi_n t_n$$

- One iteration at step τ
 - ① Choose a training pair (\mathbf{x}_n, t_n)
 - ② Update the weight vector \mathbf{w} by

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_P(\mathbf{w}) = \mathbf{w}^{(\tau)} + \eta \phi_n t_n$$

- As $y(\mathbf{x}, \mathbf{w})$ does not depend on the norm of \mathbf{w} , one can set $\eta = 1$



$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \phi_n t_n$$

The Perceptron Algorithm - Update 1



Update of the perceptron weights from a misclassified pattern
(green)

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \phi_n t_n$$





The Perceptron Algorithm - Update 2



Update of the perceptron weights from a misclassified pattern
(green)

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \phi_n t_n$$



The Perceptron Algorithm - Convergence



Classification

Generalised Linear
Model

Inference and Decision

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Does the algorithm converge ?
- For a single update step

$$-\mathbf{w}^{(\tau+1)T} \phi_n t_n = -\mathbf{w}^{(\tau)T} \phi_n t_n - (\phi_n t_n)^T \phi_n t_n < -\mathbf{w}^{(\tau)T} \phi_n t_n$$

because $(\phi_n t_n)^T \phi_n t_n = \|\phi_n t_n\|^2 > 0$.

- BUT: contributions to the error from the other misclassified patterns might have increased.
- AND: some correctly classified patterns might now be misclassified.
- **Perceptron Convergence Theorem** : If the training set is linearly separable, the perceptron algorithm is guaranteed to find a solution in a finite number of steps.